

OpenRLHF: An Easy-to-use, Scalable and High-performance RLHF Framework

Jian Hu^{*}, Xibin Wu, Wei Shen, Jason Klein Liu[†], Zilin Zhu
Weixun Wang, Songlin Jiang, Haoran Wang, Hao Chen,
Bin Chen, Weikai Fang, Xianyu, Yu Cao, Haotian Xu, Yiming Liu

 Team
janhu9527@gmail.com

Abstract

Large Language Models (LLMs) fine-tuned via Reinforcement Learning from Human Feedback (RLHF) and Reinforcement Learning with Verifiable Rewards (RLVR) significantly improve the alignment of human-AI values and further raise the upper bound of AI capabilities, particularly in reasoning-intensive, long-context Chain-of-Thought (long-CoT) tasks. However, existing RLHF (or RLVR) frameworks commonly face challenges such as inference bottlenecks and complexity barriers, restricting their accessibility for newcomers. To bridge this gap, we introduce **OpenRLHF**, a user-friendly, scalable, and easy-to-learn open-source RLHF framework built upon Ray, vLLM, DeepSpeed, and HuggingFace Transformers, featuring a simplified design, clear code structure, and comprehensive documentation to facilitate entry for researchers and practitioners. Experimental results show that OpenRLHF achieves superior training efficiency with speedups ranging from $1.22\times$ to $1.68\times$ across different model sizes compared to state-of-the-art frameworks, while requiring significantly fewer lines of code for implementation. OpenRLHF is publicly available at <https://github.com/OpenRLHF/OpenRLHF>, and has already been adopted by leading institutions to accelerate RLHF research and learning.

1 Introduction

Large Language Models (LLMs) fine-tuned via **Reinforcement Learning from Human Feedback (RLHF)** and **Reinforcement Learning with Verifiable Rewards (RLVR)** have markedly advanced human-AI alignment and elevated the upper bound of AI capabilities (Christiano et al., 2017; Stiennon et al., 2020; Guo et al., 2025; Shen et al., 2025). These approaches enable models to better conform to human intentions and values while achieving

superior reasoning performance. Notably, models such as GPT-4 (Achiam et al., 2023), DeepSeek-R1 (Guo et al., 2025), and Claude (Askell et al., 2021) excel at complex reasoning tasks by generating detailed step-by-step rationales, commonly referred to as *Chain-of-Thought* (CoT) outputs.

However, RLHF and RLVR training methodologies—especially those employing Proximal Policy Optimization (PPO)—face significant computational challenges. In particular, the inference phase often accounts for over 90% of the total RLHF (or RLVR) runtime, as models need to generate thousands of tokens during each inference step. Consequently, there is an increasing demand for efficient and scalable frameworks that reduce inference overhead and simplify distributed RLHF and RLVR training workflows.

Existing RLHF and RLVR systems such as DeepSpeed-Chat (Yao et al., 2023), TRL (von Werra et al., 2020), and ColossalChat (Li et al., 2023) have made noteworthy progress in distributed computation and memory efficiency. Conversely, industrial-grade solutions like Nemo-aligner (Shen et al., 2024a) and ChatLearn (alibaba, 2017) offer advanced optimizations at the modeling and framework levels. While Verl (Sheng et al., 2024), a framework proposed after our initial development, also provides sophisticated optimizations (e.g., its 3D-Hybrid engine), these industrial solutions generally feature tightly coupled and specialized designs that introduce considerable complexity, steep learning curves, and accessibility barriers for newcomers and academic researchers. Yet, their tightly coupled and specialized designs introduce considerable complexity, steep learning curves, and accessibility barriers for newcomers and academic researchers. Therefore, there remains a pressing need for an RLHF and RLVR framework that balances high performance, scalability, and ease of use—one that is straightforward enough for researchers new to the field, yet adaptable to

^{*}Full contributor list available in Appendix C.

[†]Ubiquant

diverse and evolving workloads.

In this paper, we present **OpenRLHF**, a simple, high-performance, fully open-source framework supporting both RLHF and RLVR, built upon Ray (Liang et al., 2018), vLLM, DeepSpeed (Yao et al., 2023), and HuggingFace Transformers (Wolf et al., 2020). OpenRLHF offers three key contributions:

- **First Ray-Based Open-Source RLHF and RLVR Architecture:** Leveraging Ray’s flexible distributed computing primitives, OpenRLHF enables streamlined orchestration and resource management for RLHF and RLVR workflows, significantly simplifying distributed operations and deployments while enhancing usability and flexibility.
- **3D Parallelism with DeepSpeed-ZeRO and Ring Attention:** To enable seamless and efficient scalability for large models, OpenRLHF integrates automatic tensor parallelism (AutoTP) provided by DeepSpeed-ZeRO and implements sequence parallelism via ring attention. This streamlined integration realizes an efficient "3D" parallel strategy—combining tensor, data, and sequence parallelism—without requiring complex engineering or extensive user configuration.
- **First Accelerated CoT Inference with vLLM:** Addressing the critical inference bottleneck in long-chain-of-thought RLHF and RLVR workloads, OpenRLHF incorporates the state-of-the-art vLLM inference engine. This integration accelerates inference through token-level parallel decoding, advanced caching, and optimized dynamic batching, thereby substantially improving overall training runtime efficiency.
- **Asynchronous Dataflow and Remote Engine Interactions:** OpenRLHF supports asynchronous dataflow and remote engine communication to maximize system throughput and resource utilization during distributed RLHF training. In this architecture, rollout engines, actor engines, and remote engines operate independently and communicate via message passing, enabling immediate processing as soon as data becomes available. This design reduces idle time, improves pipeline efficiency, and enhances scalability and flexibility across large distributed GPU clusters. By leveraging asynchronous remote engine interactions, OpenRLHF can be easily extended to support scalable **agent RL** training.

Together, these innovations position OpenRLHF as an accessible yet powerful framework suitable for both efficient experimentation in academic environments and practical deployment scenarios. Already adopted by leading institutions and companies—including CMU, MIT, Microsoft, and HKUST—OpenRLHF demonstrates broad applicability and impact across the research community (see Appendix A for details on the framework’s broad impact and adoption). We publicly release OpenRLHF to foster openness, accelerate research, and promote innovation within the RLHF and RLVR ecosystem.

2 Related Work

Reinforcement Learning from Human Feedback and Reinforcement Learning with Verifiable Rewards. Reinforcement Learning from Human Feedback (RLHF) has emerged as a pivotal paradigm for aligning large language models with human preferences (Christiano et al., 2017; Stienon et al., 2020). The foundational RLHF framework trains a reward model from human preference data and optimizes the language model using reinforcement learning algorithms such as PPO (Schulman et al., 2017). This approach has been successfully deployed in prominent models including InstructGPT (Ouyang et al., 2022), ChatGPT, and GPT-4 (Achiam et al., 2023). Building upon RLHF, Reinforcement Learning with Verifiable Rewards (RLVR) leverages automatically verifiable signals from mathematical verification, code execution, or other automated evaluation mechanisms (Guo et al., 2025; Shen et al., 2025; Cobbe et al., 2021). While PPO remains dominant for its stability (Schulman et al., 2017; Ziegler et al., 2019), alternative approaches such as Direct Preference Optimization (DPO) (Rafailov et al., 2023) have gained attention for computational efficiency. However, both RLHF and RLVR methods require substantial computational resources and sophisticated distributed training strategies. The computational challenges have motivated various optimization techniques, including efficient inference engines (Kwon et al., 2023), memory-efficient training strategies (Rajbhandari et al., 2020), and distributed orchestration frameworks (Yao et al., 2023). Nevertheless, existing solutions often sacrifice either performance for simplicity or accessibility for optimization, creating a gap that OpenRLHF aims to address.

RLHF (RLVR) Frameworks. The computational

complexity of RLHF and RLVR training has driven specialized framework development. General-purpose RL frameworks (Liang et al., 2018; Dhariwal et al., 2017; Shen et al., 2020) designed for small-scale networks fail to address LLM-specific challenges and typically employ multi-controller architectures with complex inter-process communication, creating steep learning curves. RLHF-specific frameworks face significant performance-accessibility trade-offs. Open-source solutions like TRL (von Werra et al., 2020), DeepSpeed-Chat (Yao et al., 2023), and ColossalChat (Li et al., 2023) provide accessible implementations but often lack sophisticated orchestration capabilities and struggle with inference optimization. Industrial frameworks like Nemo-aligner (Shen et al., 2024a), ChatLearn (alibaba, 2017), and Verl (Sheng et al., 2024) offer advanced optimizations including 3D parallelism and memory management. However, they feature tightly coupled architectures requiring substantial engineering expertise and extensive infrastructure setup, creating accessibility barriers for academic researchers. These systems adopt static resource allocation paradigms, leading to suboptimal utilization and limited adaptability. Most frameworks inadequately address both inference bottlenecks and usability challenges simultaneously, forcing users to choose between high performance and ease of use.

3 Design of OpenRLHF

3.1 Overview: Ray-based RLHF architecture

OpenRLHF is the first open-source, Ray-based RLHF architecture that assigns a batch of GPUs to distinct roles and manages both data flow and workflow among these roles using Ray’s scheduling capabilities. As illustrated in Figure 1, it defines two primary roles: the **rollout engine**, responsible for response generation to given prompts, and the **ZeRO engine**, which computes *logprobs*, *reference policy logprobs*, and handles model training (For a detailed design of the PPO workflow, refer to Appendix B, and for in-depth implementation tips, consult Blog (Shen et al., 2024b).)

Additionally, we leverage **VLLM** as the rollout engine, enabling efficient response generation with minimal GPU memory usage. For model training, we adopt DeepSpeed, which implements 3D parallelism, including automatic tensor parallelism, ZeRO/data parallelism, and sequence parallelism, to train both the actor and value models efficiently.




Figure 1: Overall architecture of OpenRLHF. The system assigns GPUs to two primary roles: rollout engines dedicated to response generation and actor engines responsible for computing log-probabilities and model training. OpenRLHF leverages Ray for distributed scheduling, integrates VLLM to achieve efficient response rollout with low GPU memory usage, and employs DeepSpeed-Zero for 3D parallelism (including tensor, data, and sequence parallelism) to enable efficient training. The underlying models are instantiated with flexible Transformer architectures, making the system easy to extend and adapt for diverse scenarios.

Core reason for ease of use: simplified model slicing, seamless integration, and flexible scheduling. The exchange of model weights between the rollout engine and the training engine is enabled by a flexible slicing and partitioning pipeline. **Hugging Face Transformer** (Jain, 2022) models are instantiated and trained using **DeepSpeed ZeRO**, **AutoTP**, and **Ring-Attention** (Liu et al., 2023) model parallelism. These model slices are then efficiently transferred to **VLLM** through **AutoTP** and **AutoPP**, which dynamically partition the models into sub-modules. The **Ray-based scheduling mechanism** enables seamless switching between different model parallelism modes, such as hybrid engine and asynchronous training. **This streamlined workflow significantly reduces complexity, making the system highly user-friendly and easy to extend.** Compared with architectures such as DeepSpeed-Chat, Transformer Reinforcement Learning (TRL), or other mainstream frameworks, OpenRLHF supports asynchronous dataflow and remote engine interactions, significantly improving the overall efficiency of the training process and the agent workflow.

3.2 Distributed and Efficient System Design

3D Parallelism with DeepSpeed-ZeRO and Ring Attention. To enable seamless and efficient scalability for large models, OpenRLHF integrates the latest automatic tensor parallelism (AutoTP) fea-

ture from DeepSpeed-ZeRO. In many industrial-grade RLHF architectures, users previously needed to manually specify an injection policy for each transformer model, identifying the linear layers and attention outputs that required communication between data-parallel ranks. In contrast, OpenRLHF leverages DeepSpeed-ZeRO’s new capability to support automatic tensor parallelism for HuggingFace models by default: when kernel injection is not enabled and an injection policy is not provided, DeepSpeed automatically determines and applies the necessary policy at runtime. **This greatly simplifies the user experience and extends robust tensor parallelism support to a wider range of models, removing the need for complex engineering or manual configuration.**

In addition, OpenRLHF implements sequence parallelism through ring attention. Ring attention employs a ring-based communication topology, efficiently distributing attention computation for long sequences across multiple GPUs while minimizing both memory usage and communication overhead. This is especially critical for modern RLHF and RLVR workloads involving long chain-of-thought (CoT) reasoning, where conventional attention computation can become a major scalability bottleneck. By combining AutoTP, data parallelism, and ring attention-based sequence parallelism, OpenRLHF empowers large-scale, efficient, and highly usable RLHF model training on flexible GPU clusters.

Accelerated CoT Inference with vLLM. As LLMs advance in reasoning, RLHF and RLVR pipelines increasingly face inference bottlenecks, especially with long chain-of-thought (CoT) outputs. For models like OpenAI-o1 and DeepSeek-R1, **CoT generation can dominate training time, making efficient long-form inference crucial for scalability.** To address this challenge, OpenRLHF integrates the vLLM inference framework, which is specifically designed for high-throughput and memory-efficient LLM serving. vLLM provides a streamlined interface for generating RLHF samples and supporting frequent model weight updates.

The core innovation in **vLLM** is efficient management of attention key and value memory with **PagedAttention** (Kwon et al., 2023). This technique significantly reduces memory waste to less than 4%, enabling the batching of more sequences and enhancing GPU utilization and throughput. **PagedAttention** also supports efficient memory sharing for advanced sampling methods, such as

parallel sampling and beam search, reducing memory usage by up to 55% and further boosting inference efficiency. Besides **PagedAttention**, **vLLM** has several other advantages, including continuous batching of incoming requests, fast model execution with CUDA Graph, and CUDA kernels optimized with FlashAttention and FlashInference, which enable rapid and memory-efficient attention computations. It also features speculative decoding for faster inference and chunked prefill to reduce latency on long sequences. Collectively, these enhancements make **vLLM** highly effective for large-scale, long-sequence inference, especially suitable for RLHF and RLVR pipelines, where efficiency and scalability are crucial.

Asynchronous Dataflow and Remote Engine Interactions. OpenRLHF supports asynchronous dataflow and remote engine communication to maximize system throughput and resource utilization during distributed RLHF training. In this architecture, rollout engines, actor engines, and remote engines operate independently and communicate via message passing, enabling immediate processing as soon as data becomes available. Fully asynchronous execution is especially important in the chain-of-thought (CoT) era, where inference involves generating multi-step reasoning that can vary greatly in length and computational cost. In synchronous frameworks, the slowest CoT generation can bottleneck the whole pipeline and waste resources. In contrast, OpenRLHF’s asynchronous design allows each engine to operate at its own pace, ensuring hardware is utilized even for long or variable CoT tasks. This not only accelerates training and evaluation, but also enables efficient scaling and supports dynamic agent RL workflows. Leveraging asynchronous remote engine interactions, OpenRLHF is readily extensible for scalable **agent RL** training in modern, CoT-centric environments.

4 Experiments

4.1 Performance Comparison

Long CoT Experiment Setup To ensure the applicability of the compared methods in current RLHF workflows, we conduct our experimental evaluation under a long-chain-of-thought (CoT) generation scenario. Given computational resource constraints, we focus our comparison on the long CoT RLVR setting, benchmarking OpenRLHF against Verl, currently the state-of-the-art frame-

Model Size	1K		2K		4K		8K		Avg. Speedup
	OpenRLHF (sec)	Verl (sec)	OpenRLHF (sec)	Verl (sec)	OpenRLHF (sec)	Verl (sec)	OpenRLHF (sec)	Verl (sec)	
1.5B	14.9	16.2	26.5	33.1	61.6	65.5	113.0	134.8	1.22×
7B	16.0	17.3	30.3	47.3	90.3	101.3	226.4	232.4	1.42×
14B	25.5	28.5	51.0	74.3	136.3	202.8	328.6	511.1	1.68×

Table 1: The average training time (seconds) per step for different model sizes and context window lengths compared between OpenRLHF and Verl. The speedup is calculated as the geometric mean of Verl time / OpenRLHF time across all sequence lengths.

work for RLHF training. We evaluate the training efficiency of OpenRLHF (v0.8.5) and Verl (v0.4.0) by measuring the average per-step training time (in seconds) across various model sizes (1.5B, 7B, and 14B parameters) and maximum generation lengths (1K, 2K, 4K, and 8K tokens). To ensure the base models can produce sufficiently long contextual outputs for stress testing, we adopt the DeepSeek open-source distilled Qwen series. All models are fine-tuned using the Decoupled Clip and Dynamic Sampling Policy Optimization (DAPO) algorithm (Yu et al., 2025) under identical hyperparameter settings. Experiments are conducted on 8 NVIDIA H200 140GB GPUs using PyTorch 2.7 (Imambi et al., 2021) and the ZeRO Stage 3 optimizer or Fully Sharded Data Parallel (FSDP) (Zhao et al., 2023). For each configuration, the local batch size is set to 1 to mitigate the risk of out-of-memory errors, with a maximum input context length of 1024 tokens. The reported values represent the average training time per step, excluding the first 10 steps.

Long CoT Performance Analysis The experimental results in Table 1 show that OpenRLHF consistently achieves superior training speed compared to Verl across all configurations. OpenRLHF delivers speedups ranging from 1.22× for the 1.5B model to 1.68× for the 14B model, with performance advantages becoming more pronounced as model size and context length increase. For instance, in the 14B-8K setting, OpenRLHF achieves a 1.56× speedup (328.6 seconds vs. 511.1 seconds), while the 7B-2K configuration shows a 1.56× speedup (30.3 seconds vs. 47.3 seconds). These speedup improvements can be attributed to OpenRLHF’s algorithmic design, including the DAPO optimization strategy, which effectively mitigates memory overhead and computational bottlenecks under long-context scenarios. The consistent speedup across different scales underscores OpenRLHF’s efficiency advantages in contemporary RLHF pipelines.

General RLVR Experiment To ensure a fair and controlled evaluation of training efficiency in reinforcement learning fine-tuning, we compare OpenRLHF with the optimized TRL framework on the GSM8K dataset (Cobbe et al., 2021) using the GRPO algorithm over a single training epoch. Both systems are configured with identical hyperparameters and run on the same hardware setup to isolate the effect of framework design on performance. The GRPO algorithm is selected due to its relevance in reward-based fine-tuning scenarios, and GSM8K serves as a representative benchmark for arithmetic reasoning tasks. In terms of training efficiency, OpenRLHF demonstrates a substantial advantage, completing one epoch in 1,657 seconds, compared to the 5,189 seconds required by TRL, representing approximately a 3.1× speedup. The result highlights the superior efficiency and maintainability of OpenRLHF’s implementation, which benefits from a streamlined design and targeted system-level optimizations.

General RLHF Experiment To evaluate the training efficiency of modern RLHF frameworks, we compare OpenRLHF with the optimized DeepSpeed-Chat (DSChat) implementation. The experiment involves fine-tuning 1,024 prompts using the PPO algorithm for one epoch under identical hardware and hyperparameter configurations. This setup ensures that observed differences in performance are attributable solely to differences in system design and optimization strategies between the two frameworks. In terms of training time, OpenRLHF completes the task in 236.8 seconds, significantly outperforming DSChat, which requires 855.09 seconds, resulting in a 3.6× speedup. This performance gain is primarily driven by two system-level innovations in OpenRLHF: the use of vLLM for accelerated token generation and Ray for efficient distributed execution. Collectively, these design choices result in a more scalable and excellent RLHF framework.




Figure 2: Core code complexity comparison across RLHF frameworks (lines of code). Lower values indicate simpler implementation and better maintainability.

4.2 Usability Comparison

As illustrated in Figure 2, OpenRLHF achieves a compelling balance between implementation conciseness and training performance. Despite being the second most concise framework with only 8,523 lines of code—significantly fewer than TRL (19,071) and Verl (32,325)—OpenRLHF demonstrates a clear performance advantage across standard RLHF benchmarks. This streamlined codebase not only facilitates easier comprehension and modification for developers but also reduces the engineering overhead associated with integration into custom pipelines. In addition to its lightweight design, OpenRLHF offers comprehensive support for various reinforcement learning fine-tuning paradigms, including Supervised Fine-Tuning (SFT), Direct Preference Optimization (DPO) (Rafailov et al., 2023), Reward Model (RM), and Process Reward Model (PRM) (Lightman et al., 2023). This broad functionality, combined with its modular and well-documented architecture, significantly lowers the barrier to entry for both research and production use. Overall, OpenRLHF’s design exemplifies high usability by combining minimal code complexity with extensive functionality and competitive performance.

5 Limitations

While OpenRLHF demonstrates significant advantages in balancing performance and accessibility, several limitations should be acknowledged. Despite our optimization efforts, OpenRLHF may not match the peak performance of highly specialized industrial frameworks that benefit from dedicated engineering teams and extensive resources. As a community-driven open-source

project without dedicated economic support, OpenRLHF faces resource constraints in rapidly integrating cutting-edge features, potentially resulting in delays compared to well-funded commercial frameworks. Currently, the framework focuses primarily on language models and does not support Vision-Language Models or other multimodal architectures, limiting its applicability to multimodal AI alignment research. Additionally, OpenRLHF’s modular design introduces dependencies on external systems such as Ray, vLLM, and DeepSpeed, where updates in these upstream systems may require maintenance work or introduce compatibility issues. Despite these limitations, we believe OpenRLHF’s contributions to accessibility and democratization of RLHF research provide significant value to the community.

6 Conclusion

We presented OpenRLHF, a simple yet high-performance open-source framework that bridges the gap between performance and usability in RLHF and RLVR training. By integrating Ray’s distributed computing, vLLM’s inference optimization, DeepSpeed ZeRO’s memory efficiency, and ring attention’s sequence parallelism, OpenRLHF delivers four key innovations: Ray-based architecture for streamlined orchestration, 3D parallelism for efficient scaling, accelerated COT inference addressing the critical bottleneck, and asynchronous dataflow for maximum throughput. The framework’s broad adoption across leading institutions and companies—from academic courses at CMU to production deployments at major tech companies—validates its real-world effectiveness. OpenRLHF’s influence on subsequent frameworks and its role in democratizing RLHF research demonstrate its significant contribution to the field. By open-sourcing this framework, we aim to foster innovation, accelerate research progress, and enable broader participation in aligned AI development.

References

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, and 1 others. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- alibaba. 2017. chatlearn. <https://github.com/alibaba/ChatLearn>.
- Amanda Askell, Yuntao Bai, Anna Chen, Dawn Drain, Deep Ganguli, Tom Henighan, Andy Jones, Nicholas Joseph, Ben Mann, Nova DasSarma, and 1 others. 2021. A general language assistant as a laboratory for alignment. *arXiv preprint arXiv:2112.00861*.
- Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, 30.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, and 1 others. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Prafulla Dhariwal, Christopher Hesse, Oleg Klimov, Alex Nichol, Matthias Plappert, Alec Radford, John Schulman, Szymon Sidor, Yuhuai Wu, and Peter Zhokhov. 2017. Openai baselines. <https://github.com/openai/baselines>.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shitong Ma, Peiyi Wang, Xiao Bi, and 1 others. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*.
- Sagar Imambi, Kolla Bhanu Prakash, and GR Kana-gachidambaresan. 2021. Pytorch. *Programming with TensorFlow: solution for edge computing applications*, pages 87–104.
- Shashank Mohan Jain. 2022. Hugging face. In *Introduction to transformers for NLP: With the hugging face library and models to solve problems*, pages 51–67. Springer.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*.
- Shenggui Li, Hongxin Liu, Zhengda Bian, Jiarui Fang, Haichen Huang, Yuliang Liu, Boxiang Wang, and Yang You. 2023. Colossal-ai: A unified deep learning system for large-scale parallel training. In *Proceedings of the 52nd International Conference on Parallel Processing*, pages 766–775.
- Eric Liang, Richard Liaw, Robert Nishihara, Philipp Moritz, Roy Fox, Ken Goldberg, Joseph Gonzalez, Michael Jordan, and Ion Stoica. 2018. Rllib: Abstractions for distributed reinforcement learning. In *International conference on machine learning*, pages 3053–3062. PMLR.
- Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2023. Let’s verify step by step. In *The Twelfth International Conference on Learning Representations*.
- Hao Liu, Matei Zaharia, and Pieter Abbeel. 2023. Ring attention with blockwise transformers for near-infinite context. *arXiv preprint arXiv:2310.01889*.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, and 1 others. 2022. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36:53728–53741.
- Samyam Rajbhandari, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. 2020. Zero: Memory optimizations toward training trillion parameter models. In *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–16. IEEE.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.
- Gerald Shen, Zhilin Wang, Olivier Delalleau, Jiaqi Zeng, Yi Dong, Daniel Egert, Shengyang Sun, Jimmy Zhang, Sahil Jain, Ali Taghibakhshi, Markel Sanz Ausin, Ashwath Aithal, and Oleksii Kuchaiev. 2024a. **Nemo-aligner: Scalable toolkit for efficient model alignment.** Preprint, arXiv:2405.01481.
- Wei Shen, Xiaonan He, Chuheng Zhang, Qiang Ni, Wanchun Dou, and Yan Wang. 2020. Auxiliary-task based deep reinforcement learning for participant selection problem in mobile crowdsourcing. In *Proceedings of the 29th ACM International Conference on Information & Knowledge Management*, pages 1355–1364.
- Wei Shen, Jian Hu, Pengyu Zhao, Xiaonan He, and Lichang Chen. 2024b. Advanced tricks for training large language models with proximal policy optimization. <https://swtheking.notion.site/eb7b2d1891f44b3a84e7396d19d39e6f?v=01bcb084210149488d730064cbabc99f&pvs=74>. Notion Blog.

Wei Shen, Guanlin Liu, Zheng Wu, Ruofei Zhu, Qingping Yang, Chao Xin, Yu Yue, and Lin Yan. 2025. Exploring data scaling trends and effects in reinforcement learning from human feedback. *arXiv preprint arXiv:2503.22230*.

Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. 2024. Hybridflow: A flexible and efficient rlhf framework. *arXiv preprint arXiv:2409.19256*.

Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. 2020. Learning to summarize with human feedback. *Advances in Neural Information Processing Systems*, 33:3008–3021.

Leandro von Werra, Younes Belkada, Lewis Tunstall, Edward Beeching, Tristan Thrush, Nathan Lambert, and Shengyi Huang. 2020. Trl: Transformer reinforcement learning. <https://github.com/huggingface/trl>.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pieric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, and 3 others. 2020. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45, Online. Association for Computational Linguistics.

Zhewei Yao, Reza Yazdani Aminabadi, Olatunji Ruwase, Samyam Rajbhandari, Xiaoxia Wu, Ammar Ahmad Awan, Jeff Rasley, Minjia Zhang, Conglong Li, Connor Holmes, and 1 others. 2023. Deepspeed-chat: Easy, fast and affordable rlhf training of chatgpt-like models at all scales. *arXiv preprint arXiv:2308.01320*.

Qiyi Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, and 1 others. 2025. Dapo: An open-source llm reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*.

Yanli Zhao, Andrew Gu, Rohan Varma, Liang Luo, Chien-Chin Huang, Min Xu, Less Wright, Hamid Shojanazeri, Myle Ott, Sam Shleifer, and 1 others. 2023. Pytorch fsdp: experiences on scaling fully sharded data parallel. *arXiv preprint arXiv:2304.11277*.

Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*.

A Broad Impact and Adoption

OpenRLHF has achieved significant adoption and impact across both academic and industrial communities since its release. The framework’s balance of high performance and accessibility has made it a preferred choice for diverse applications ranging from research experimentation to production deployment.

OpenRLHF has been formally integrated into academic curricula, notably being adopted by The CMU Advanced Natural Language Processing Spring 2025 course as a core teaching framework. This integration demonstrates the framework’s educational value and accessibility for students new to RLHF. The framework has also received recognition from major technology communities, including an invitation to present at the PyTorch Expert Exchange 2025, highlighting its technical contributions to the broader deep learning ecosystem.

The framework has been widely adopted by leading technology companies and research institutions, including Google, ByteDance, Baidu, NVIDIA, Tencent, China Telecom, Vivo, Nexus-Flow, JSC, UC Berkeley’s Starling Team, Meituan, and HKUST. This diverse adoption across major tech companies, telecommunications providers, and academic institutions demonstrates OpenRLHF’s versatility and robustness in handling various scales of RLHF workloads, from research prototyping to production-scale deployments.

OpenRLHF’s modular design and extensible architecture have enabled it to serve as a foundation for specialized frameworks. Notable examples include LMM-R1 for multimodal reinforcement learning, MARTI for advanced reasoning tasks, and MM-EUREKA for multi-modal applications. These derivative frameworks demonstrate OpenRLHF’s flexibility in supporting diverse research directions and its role as a platform for innovation in the RLHF ecosystem.

The design principles and technical innovations introduced in OpenRLHF have influenced subsequent framework development in the field. Several prominent frameworks including Verl, Alibaba ROLL, SLIME, and Open-Reasoner-Zero have acknowledged OpenRLHF’s contributions in their documentation and publications, citing its distributed architecture design, Ray-based orchestration approach, and integration strategies as influential to their own development. This acknowledgment reflects OpenRLHF’s role in advancing

the state-of-the-art in RLHF system design and establishing best practices for distributed RLHF training.

Beyond direct usage and citations, OpenRLHF has contributed to democratizing RLHF research by lowering the barrier to entry for academic researchers and smaller organizations. The framework’s emphasis on ease of use without sacrificing performance has enabled broader participation in RLHF research, fostering innovation across diverse research communities that might otherwise lack the resources for complex distributed training setups.

B PPO Workflow Design

This section presents OpenRLHF’s comprehensive PPO-based RLHF training workflow, which orchestrates multiple specialized engines to efficiently handle the complex multi-stage training process (as shown in Figure 3). The OpenRLHF PPO workflow consists of four main stages executed iteratively: (1) **Rollout Generation**, where the current policy generates responses to prompts; (2) **Reward Computation**, where responses are evaluated using a trained reward model; (3) **Advantage Estimation**, where advantages and returns are calculated using GAE (Generalized Advantage Estimation); and (4) **Policy Optimization**, where the policy is updated using PPO loss. This pipeline repeats for multiple iterations until convergence.

The training begins with the **Rollout Engine** generating responses for a batch of prompts using the current policy π_θ . A batch of prompts $\{x_1, x_2, \dots, x_B\}$ is sampled from the training dataset, and the Rollout Engine, equipped with vLLM for efficient inference, generates responses $\{y_1, y_2, \dots, y_B\}$ using the current policy. During generation, the engine records action log-probabilities $\log \pi_\theta(y_i|x_i)$ and attention masks. The generated sequences (x_i, y_i) along with their metadata are collected for subsequent processing. The Rollout Engine leverages vLLM’s optimizations including continuous batching, KV-cache management, and PagedAttention to maximize throughput during this inference-heavy stage.

Once rollouts are generated, the system computes rewards and reference policy log-probabilities. The trained reward model R_ϕ evaluates each prompt-response pair to produce scalar rewards $r_i = R_\phi(x_i, y_i)$. Simultaneously, the frozen reference policy π_{ref} computes log-probabilities $\log \pi_{\text{ref}}(y_i|x_i)$ for KL regularization, and the critic




Figure 3: Overall PPO workflow of OpenRLHF.

network V_ψ estimates state values $V_\psi(x_i, y_{i,t})$ for advantage computation. These computations can be parallelized across multiple GPUs and are often batched together for efficiency.

The system then computes advantages using Generalized Advantage Estimation (GAE). First, temporal difference residuals are calculated: $\delta_t = r_t + \gamma V_\psi(s_{t+1}) - V_\psi(s_t)$. Then, advantages are computed using GAE: $A_t = \sum_{l=0}^{\infty} (\gamma\lambda)^l \delta_{t+l}$, and discounted returns are calculated as $R_t = A_t + V_\psi(s_t)$. The advantage computation incorporates KL penalty terms to prevent the policy from deviating too far from the reference policy: $r'_t = r_t - \beta \text{KL}[\pi_\theta || \pi_{\text{ref}}]$.

The final stage updates the policy using PPO's clipped objective function. The **ZeRO Engine** computes the policy loss using the clipped surrogate objective: $L^{\text{CLIP}}(\theta) = \mathbb{E}[\min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon)A_t)]$, where $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$ is the probability ratio. The critic loss is computed as $L^V(\psi) = \mathbb{E}[(V_\psi(s_t) - R_t)^2]$. The critic loss is computed as $L^V(\psi) = \mathbb{E}[(V_\psi(s_t) - R_t)^2]$. The total loss combines policy and value losses: $L_{\text{total}} = L^{\text{CLIP}} + c_1 L^V + c_2 S[\pi_\theta](s_t)$, where c_1 and c_2 are coefficients for value loss and entropy bonus respectively, and $S[\pi_\theta](s_t)$ is the entropy of the policy distribution to encourage exploration. The ZeRO Engine performs gradient computation and model parameter

updates using DeepSpeed ZeRO optimizations for memory efficiency and scalability.

Throughout this workflow, OpenRLHF's Ray-based architecture enables seamless coordination between different engines while maintaining computational efficiency. The Rollout Engine and ZeRO Engine can operate on different GPU clusters optimized for their respective workloads— inference-optimized hardware for rollout generation and training-optimized hardware for policy updates. Ray's distributed scheduling automatically handles data transfer, synchronization, and fault tolerance across the distributed system. This asynchronous execution model allows the system to overlap computation stages where possible, significantly improving overall training throughput compared to traditional synchronous RLHF implementations.

The entire PPO training loop continues for multiple epochs, with each epoch processing multiple batches of rollout data. Early stopping criteria based on KL divergence thresholds and performance metrics prevent policy collapse and ensure stable training. The modular design allows for easy integration of advanced techniques such as advantage normalization, gradient clipping, and adaptive KL penalty coefficients, making OpenRLHF highly customizable for different research and production scenarios.

C Full Contributors

- **Ray Integration:** Jian Hu, Xibin Wu, Songlin Jiang
- **vLLM Integration:** Jian Hu, Xibin Wu, Songlin Jiang, Kaichao You (vLLM Team)
- **Ring Attention:** Zilin Zhu (Zhipu), Zhibo Zhou (Vivo), gzpan(GitHub User), Jian Hu
- **Supervised Learning:** Haoran Wang and Xianyu
- **DeepSpeed and Transformers Integration:** Jian Hu, Xibin Wu, Songlin Jiang, Bin Chen, Yiming Liu
- **Asynchronous Agentic RL:** Haotian Xu and Jian Hu
- **Single Controller Architecture:** Jian Hu
- **PPO Implementation:** Jian Hu
- **GRPO Implementation:** Jason Klein Liu
- **KL Control Mechanism:** Yiming Liu and Jason Klein Liu
- **RLVR Experiment Design:** Jason Klein Liu
- **RLHF Experiment Design:** Jian Hu
- **Dynamic Sampling:** Jian Hu
- **Dynamic Batching:** Hao Chen
- **Documentation:** Jian Hu
- **Testing and Bug Reports:** Weikai Fang, Wei Shen, Weixun Wang
- **Paper Writing and Presentations:** Wei Shen, Jason Klein Liu, Weixun Wang, Yu Cao, Jian Hu

A more complete list can be found in the Open-RLHF commit and release history.